

Applied Economic Forecasting

7. ARIMA Models

1 Stationarity and Differencing

2 AR Models

3 MA Models

4 ARIMA Modeling

5 Nonseasonal ARIMA Models

6 Seasonal ARIMA Models

Section 1

Stationarity and Differencing

Stationarity

Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

A **stationary series** is:

- constant mean
- constant variance
- no patterns predictable in the long-term
- No Trend
- No Seasonality

A Note on Cyclic Series


A time series with cyclic behavior (but with no trend or seasonality) is stationary. This is because the cycles are not of a fixed length, so we are unable to predict where the peaks and troughs of the cycles will be.

Stationary?


Google Closing Price




Change in Google Price




AUS: Beer production




AUS: Gas production




CAN: Lynx Trapping



AUS: Pigs Slaughtered



Stationary?



Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

- The transformations from earlier modules help to **stabilize the variance**.
- For ARIMA modeling, we also need to **stabilize the mean**.

Non-stationarity in the mean

Identifying non-stationary series

- time plot.
- The ACF of stationary data drops to zero relatively quickly.
- The ACF of non-stationary data decreases slowly.
 - Think back to the ACF of a series with a trend.
- For non-stationary data, the value of r_1 (the first lag of the ACF) is often large and positive.

Non-stationarity in the mean: Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series:

$$\Delta y_t = y_t - y_{t-1}.$$

- The differenced series will have only $T - 1$ values since it is not possible to calculate a difference, Δy_1 , for the first observation.


Non-stationarity in the mean: Differencing

```
goog <- gafa_stock %>% filter(Symbol == "GOOG") %>%  
  #First Diff  
  mutate(dg = difference(Close))  
  
p1 <- goog %>% autoplot(Close) + ggtitle("Google Price")  
p2 <- goog %>% autoplot(dg, col = "blue4") +  
  ggtitle(bquote(Delta~"Price"))  
  
## ACFs  
p3 <- goog %>% ACF(Close) %>% autoplot() +  
  ggtitle("ACF: Price")  
p4 <- goog %>% ACF(dg) %>% autoplot() +  
  ggtitle(bquote("ACF:"~Delta~"Price"))  
  
gridExtra::grid.arrange(p1,p3,p2,p4, ncol = 2)
```


Google Price




ACF: Price



Δ Price



ACF: Δ Price



Non-stationarity in the mean: Differencing

Random Walk Model

When the differenced series is white noise, the model for the original series can be written as:

$$y_t - y_{t-1} = \varepsilon_t$$

where ε_t denotes white noise. Rearranging this leads to the “random walk” model:

$$y_t = y_{t-1} + \varepsilon_t$$

Random walk models are widely used for non-stationary data, particularly financial and economic data. Random walks typically have:

- long periods of apparent trends up or down
- sudden and unpredictable changes in direction.

Non-stationarity in the mean: Differencing

Random Walk w. Drift

A closely related model allows the differences to have a non-zero mean. Then

$$y_t = c + y_{t-1} + \varepsilon_t \quad \text{OR} \quad y_t - y_{t-1} = c + \varepsilon_t$$

c represents the average of the changes between consecutive observations.

- c is positive, then the average change is an increase in the value of y_t . Thus, y_t will tend to drift upwards.
- If c is negative, y_t will tend to drift downwards.

Non-stationarity in the mean: Differencing

Second-order differencing

Occasionally, the differenced data will not appear stationary and it may be necessary to difference the data a second time:

$$\begin{aligned}\Delta^2 y_t &= \Delta y_t - \Delta y_{t-1} \\ &= (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) \\ &= y_t - 2y_{t-1} + y_{t-2}.\end{aligned}$$

- $\Delta^2 y_t$ will have $T - 2$ values.
- In practice, it is almost never necessary to go beyond second-order differences.

Seasonal differencing

A seasonal difference is the difference between an observation and the corresponding observation **from that season in the previous year**.

$$\Delta_m y_t = y_t - y_{t-m}$$


where m = number of seasons.

- For monthly data $m = 12$.
- For quarterly data $m = 4$.

Seasonal differencing


Stable Variance?

```
a10 <- PBS%>%  
  filter(ATC2 == "A10") %>%  
  summarise(Cost = sum(Cost)/1e6)  
  
a10 %>% autoplot(Cost, col = "purple") + ggtitle("Antidiabetic drug sales")
```




Seasonal differencing

```
a10 %>% autoplot(Cost %>% log(), col = "green4") +  
  ggtitle("Logged Antidiabetic drug sales")
```



Seasonal differencing

```
a10 %>% autoplot(Cost %>% log() %>% difference(12),  
                  col = "darkblue") +  
  labs(title = "Logged Antidiabetic drug sales",  
        subtitle = "Seasonally differenced")
```



Seasonal differencing

- Seasonally differenced series is closer to being stationary.
- Remaining non-stationarity can be removed with further first difference.


If $\Delta_m y_t = y_t - y_{t-12}$ denotes seasonally differenced series, then a twice-differenced series takes the form:

$$\begin{aligned} y_t^* &= \Delta_m y_t - \Delta_m^2 y_{t-1} \\ &= (y_t - y_{t-12}) - (y_{t-1} - y_{t-13}) \\ &= y_t - y_{t-1} - y_{t-12} + y_{t-13} . \end{aligned}$$

Generic Form $y_t^* = y_t - y_{t-1} - y_{t-m} + y_{t-m-1}$

Seasonal differencing

```
a10 %>% autoplot(difference(log(Cost), 12), col = "darkblue")  
  labs(title = "Logged Antidiabetic drug sales",  
       subtitle = "Seasonally differenced")
```



When both seasonal and first differences are applied...

- it makes no difference which is done first– the result will be the same.
- If seasonality is strong, it is recommended that seasonal differencing be done first because sometimes the resulting series will be stationary and there will be no need for further first difference.

It is important that if differencing is used, the differences are interpretable.

Interpretation of differencing

- first differences are the change between **one observation and the next**;
- seasonal differences are the change between **one year to the next**.

But taking lag 3 differences for yearly data, for example, results in a model which cannot be sensibly interpreted.

Statistical tests to determine the required order of differencing.

- ① Augmented Dickey Fuller (ADF) test **Null hypothesis is that the data are non-stationary and non-seasonal.**
- ② Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test **Null hypothesis is that the data are stationary and non-seasonal.**
- ③ Other tests available for seasonal data.

ADF test


```
goog %>% features(Close, unitroot_kpss)
```

```
## # A tibble: 1 x 3  
##   Symbol kpss_stat kpss_pvalue  
##   <chr>      <dbl>      <dbl>  
## 1 GOOG      14.8        0.01
```

Takeaway: The p-value of 0.01 suggest that we can reject the null (of stationarity) at the 1% significance level. This implies that the variable is actually nonstationary and will need to be differenced.

Note

- By default, p-values smaller than 0.01 will be reported as 0.01 exactly.
- Those p-values larger than 0.1 will be reported as 0.1 exactly.

```
goog %>% features(Close %>% difference(), unitroot_kpss)
```

```
## # A tibble: 1 x 3  
##   Symbol kpss_stat kpss_pvalue  
##   <chr>      <dbl>      <dbl>  
## 1 GOOG      0.0529      0.1
```

Takeaways: The pvalue of the KPSS test is rather large. Since it is greater than 10%, we fail to reject the null of stationarity and conclude that the differenced series does appear to be stationary.

KPSS test

Determine the optimal number of times to difference using package.

Non-Seasonal Differencing:

```
goog %>% features(Close, unitroot_ndiffs)
```

```
## # A tibble: 1 x 2
##   Symbol ndiffs
##   <chr>   <int>
## 1 GOOG      1
```

Seasonal Differencing:

```
a10 %>% features(Cost %>% log(), unitroot_nsdiffs)
```

```
## # A tibble: 1 x 1
##   nsdiffs
##   <int>
## 1      1
```

Was it necessary to doubly difference the a10 series earlier?

```
a10 <- a10 %>% mutate(l.cost = log(Cost))  
# Seasonal Difference  
a10 %>% features(l.cost, unitroot_nsdiffs)
```

```
## # A tibble: 1 x 1  
##   nsdiffs  
##   <int>  
## 1      1
```

```
# First Difference of the seasonally differenced series  
a10 %>% features(l.cost %>% difference(12), unitroot_nsdiffs)
```

```
## # A tibble: 1 x 1  
##   ndiffs  
##   <int>  
## 1      0
```

Your turn!

Using the codes below, determine the appropriate differencing, **after variance stabilization, if necessary**, that must be done to the Turnover series to render it stationary.

```
aus_retail <- aus_retail %>%  
  summarise(Turnover = sum(Turnover))
```

Section 2

AR Models

In an autoregression model, we forecast our variable of interest (y_t) using a linear combination of **past values** of that variable.

Therefore, the term autoregression indicates that it is a regression of the variable against itself.

An autoregressive model of lag order p can be expressed as:

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \varepsilon_t$$


where ε_t is white noise.

$$y_t = c + \phi_1 y_{t-1} + \varepsilon_t$$


- When $\phi_1 = 0$, y_t is **equivalent to WN**
- When $\phi_1 = 1$ and $c = 0$, y_t is **equivalent to a RW**
- When $\phi_1 = 1$ and $c \neq 0$, y_t is **equivalent to a RW with drift**
- When $\phi_1 < 0$, y_t tends to **oscillate between positive and negative values**.

Sample of an AR(1) Process


AR(1) $\phi = +0.9$



AR(1) $\phi = -0.9$



Sample of an AR(2) Process



Stationarity conditions

We normally restrict autoregressive models to stationary data, and then some constraints on the values of the parameters are required.

General condition for stationarity

Complex roots of

$$1 - \phi_1 z - \phi_2 z^2 - \dots - \phi_p z^p$$

lie outside the unit circle on the complex plane.

- For $p = 1$: $-1 < \phi_1 < 1$.
- For $p = 2$:

$$-1 < \phi_2 < 1, \quad \phi_2 + \phi_1 < 1, \quad \phi_2 - \phi_1 < 1.$$

- More complicated conditions hold for $p \geq 3$.
- Estimation software takes care of this.

Section 3

MA Models

Moving Averages (MA)

Moving Average (MA) models:

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$


where ε_t is white noise.

This is a multiple regression with ***past errors*** as predictors.


Do not confuse this with moving average smoothing!

Sample MA Models

$$\text{MA}(1): y_t = 20 + \varepsilon_t + 0.8\varepsilon_{t-1}$$



$$\text{MA}(2): y_t = \varepsilon_t - \varepsilon_{t-1} + 0.8\varepsilon_{t-2}$$



It is possible to write any stationary AR(p) process as an MA(∞) process.

Example: AR(1)

$$\begin{aligned}y_t &= \phi_1 y_{t-1} + \varepsilon_t \\&= \phi_1 (\phi_1 y_{t-2} + \varepsilon_{t-1}) + \varepsilon_t \\&= \phi_1^2 y_{t-2} + \phi_1 \varepsilon_{t-1} + \varepsilon_t \\&= \phi_1^3 y_{t-3} + \phi_1^2 \varepsilon_{t-2} + \phi_1 \varepsilon_{t-1} + \varepsilon_t \\&\dots\end{aligned}$$

$$y_t = \phi^k y_{t-k} + \sum_{k=0}^{\infty} \phi^k \varepsilon_{t-k}.$$

Provided $-1 < \phi_1 < 1$:

$$y_t = \varepsilon_t + \phi_1 \varepsilon_{t-1} + \phi_1^2 \varepsilon_{t-2} + \phi_1^3 \varepsilon_{t-3} + \dots$$

Invertibility

- Any $MA(q)$ process can be written as an $AR(\infty)$ process if we impose some constraints on the MA parameters.
- Then the MA model is called “invertible”.
- Invertible models have some mathematical properties that make them easier to use in practice.
- Invertibility of an ARIMA model is equivalent to forecastability of an ETS model.

General condition for invertibility

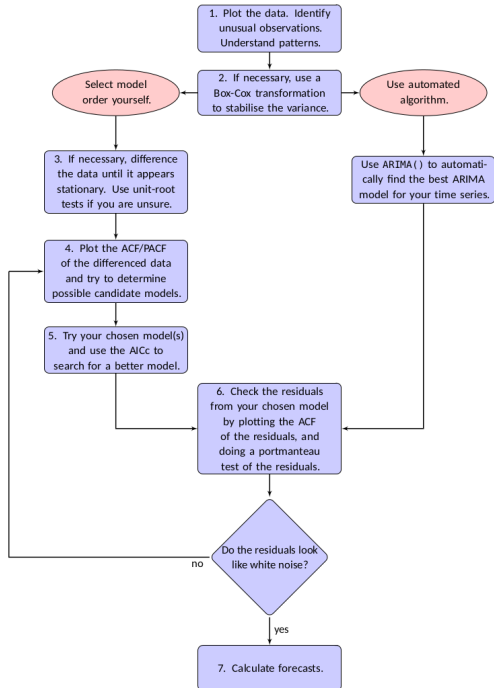
Complex roots of $1 + \theta_1 z + \theta_2 z^2 + \dots + \theta_q z^q$ lie outside the unit circle on the complex plane.

- For $q = 1$: $-1 < \theta_1 < 1$.
- For $q = 2$:
 $-1 < \theta_2 < 1 \quad \theta_2 + \theta_1 > -1 \quad \theta_1 - \theta_2 < 1$.
- More complicated conditions hold for $q \geq 3$.
- Estimation software takes care of this.

Section 4

ARIMA Modeling

The Box-Jenkins methodology refers to a set of procedures for identifying, fitting, and checking ARIMA models with time series data. Forecasts follow directly from the form of the fitted model.



- Autoregressive integrated moving average (ARIMA) models
 - a class of linear models that are capable of representing *stationary* as well as *nonstationary* time series.
 - it generates forecasts based on a description of *historical patterns in the data itself*,

$$Y_t = \beta_0 + \underbrace{\beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \cdots + \beta_k Y_{t-p}}_{\text{autoregressive terms}} \\ + \underbrace{\varepsilon_t - \omega_1 \varepsilon_{t-1} - \omega_2 \varepsilon_{t-2} - \cdots - \omega_q \varepsilon_{t-q}}_{\text{moving average terms}}$$

Nonstationary time series data

- ① The mean of the data is function of time

$$E(Y_t) = a + b \cdot t$$

so-called first moment nonstationary.

- A unit root (random walk)

$$Y_t = Y_{t-1} + \varepsilon_t$$

Stationary time series data

- The mean of the data is independent of time (i.e., constant),

$$E(Y_t) = a$$

so-called first moment stationary.

- The initial selection of an ARIMA model is based on
 - ① Autocorrelation function (ACF) ==> **Will help us to determine the moving average (MA) component.**

$$r_k = \frac{\text{cov}(Y_t, Y_{t-k})}{\sqrt{\text{var}(Y_t)}\sqrt{\text{var}(Y_{t-k})}}$$

- ② Partial autocorrelation function (PACF) ==> **Will help us to determine the autoregressive (AR) component.**

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \cdots + \beta_k Y_{t-k} + \varepsilon_t$$

such that

$$r_k = \beta_k$$

Box-Jenkins Methodology

Distance k	Regression	Partial Autocorrelation coefficient r_k
1	$Y_t = \beta_0 + \beta_1 Y_{t-1} + \varepsilon_t$	β_1
2	$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \varepsilon_t$	β_2
3	$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \beta_3 Y_{t-3} + \varepsilon_t$	β_3
...
k	$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \cdots + \beta_k Y_{t-k} + \varepsilon_t$	β_k

Section 5

Nonseasonal ARIMA Models

Identifying Nonseasonal AR and MA Models

	<i>Autocorrelations</i>	<i>Partial Autocorrelations</i>
MA(q)	Cut off after the order q of the process	Die out
AR(p)	Die out	Cut off after the order p of the process
ARMA(p, q)	Die out	Die out

LET US EXPLORE SOME SIMULATIONS AND TEST OUR THEORY!

Autoregressive Moving Average models:

$$y_t = c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} \\ + \theta_1 \varepsilon_{t-1} + \cdots + \theta_q \varepsilon_{t-q} + \varepsilon_t.$$

- Predictors include both **lagged values of y_t** and **lagged errors**.
- Conditions on coefficients ensure stationarity.
- Conditions on coefficients ensure invertibility.

Autoregressive Integrated Moving Average models

- Combine ARMA model with **differencing**.

Nonseasonal ARIMA models

Autoregressive Integrated Moving Average models

ARIMA(p, d, q) model

- AR: p = order of the autoregressive part
I: d = degree of first differencing involved
MA: q = order of the moving average part.


Special cases of Nonseasonal ARIMA Models

- White noise model: ARIMA(0,0,0) with no constant
- Random walk: ARIMA(0,1,0) with no constant
- Random walk with drift: ARIMA(0,1,0) with constant
- AR(p): ARIMA($p,0,0$)
- MA(q): ARIMA(0,0, q)

Step (1): Plot Data

```
egy <- global_economy |>  
  filter(Code == "EGY")  
  
egy |> autoplot(Exports) +  
  labs(y = "% of GDP", title = "Egyptian exports")
```

Egyptian exports



Step 2: Stabilize Variance?

Step 3: Check if Stationary

```
egy |> features(Exports, unitroot_kpss)
```

```
## # A tibble: 1 x 3
##   Country          kpss_stat kpss_pvalue
##   <fct>           <dbl>     <dbl>
## 1 Egypt, Arab Rep.    0.192       0.1
```

Key Takeaways:


- The large p-value of the KPSS test suggests that we **fail to reject** the null of stationarity at all conventional levels of significance.
- Our unit root test therefore, suggests that differencing is not necessary.
 - If you were not able to see this at the outset, look again and notice that the series seems to represent somewhat of a cycle in Exports.
Recall from earlier, that cycles are stationary!

ARIMA in Motion

Step 4: Plot ACF and PACFs to determine potential model candidates.

```
egy |> gg_tsdisplay(Exports, plot_type = c("partial")) +  
  labs(title = "ACF and PACF Plots of Egyptian Exports")
```

ACF and PACF Plots of Egyptian Exports



Takeaways

The ACF decays in a 'sinusoidal' pattern while the last significant lag in the PACF is lag 4. We could explore the possibility of an AR(4) model. For this exercise, we will assume that there is only 1 candidate model.

```
mod.fit <- egypt > model(ARIMA(Exports ~ pdq(4,0,0)))  
mod.fit > report()
```

```
## Series: Exports  
## Model: ARIMA(4,0,0) w/ mean  
##  
## Coefficients:  
##          ar1      ar2      ar3      ar4  constant  
##      0.9861 -0.1715  0.1807 -0.3283    6.6922  
## s.e.  0.1247   0.1865  0.1865   0.1273    0.3562  
##  
## sigma^2 estimated as 7.885:  log likelihood=-140.5  
## AIC=293.1   AICc=294.7   BIC=305.4
```


What about other potential models? How about the best ARIMA models with $p \in \{1, 2, 4\}$ and $q \in \{0, 1, 2, 3\}$


```
egy |> model(ARIMA(Exports ~ pdq(p = c(1,2,4),  
                                q = 0:3))) |>  
  report()
```

```
## Series: Exports  
## Model: ARIMA(2,0,1) w/ mean  
##  
## Coefficients:  
##           ar1      ar2      ma1  constant  
##          1.6764 -0.8034 -0.6896    2.5623  
## s.e.    0.1111   0.0928   0.1492    0.1161  
##  
## sigma^2 estimated as 8.046:  log likelihood=-141.6  
## AIC=293.1   AICc=294.3   BIC=303.4
```

ARIMA in Motion

Step 6: Check residuals

```
mod.fit |> gg_tsresiduals()
```



Step 6: Check residuals


```
augment(mod.fit) |>
  features(.innov, ljung_box, lag = 10, dof = 4)
#where dof (K) = p+q, is the number of AR & MA parameters in model
```

```
## # A tibble: 1 x 4
##   Country      .model                                lb_stat lb_pvalue
##   <fct>        <chr>                                <dbl>    <dbl>
## 1 Egypt, Arab Rep. ARIMA(Exports ~ pdq(4, 0, 0))      3.95      0.684
```

ARIMA in Motion

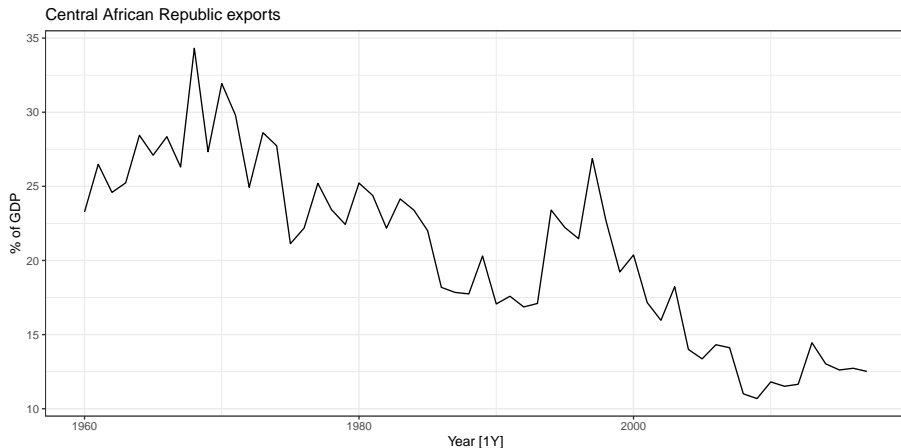
Step 7: Forecast using the preferred Model

```
mod.fit |> forecast(h = 10) |>  
  autoplot(egy) +  
  labs(title = "Egyptian Exports Forecasts",  
        subtitle = "ARIMA(4,0,0)")
```



Example #2:

```
caf <- global_economy |> filter(Code == "CAF")  
  
caf |> autoplot(Exports) +  
  labs(title="Central African Republic exports", y="% of GDP")
```



```
caf |> features(Exports, unitroot_kpss)
```

```
## # A tibble: 1 x 3
```

##	Country	kpss_stat	kpss_pvalue
##	<fct>	<dbl>	<dbl>
## 1	Central African Republic	1.28	0.01


```
caf |> features(Exports, unitroot_ndiffs)
```

```
## # A tibble: 1 x 2
```

##	Country	ndiffs
##	<fct>	<int>
## 1	Central African Republic	1

```
caf <- caf |> mutate(d.exports = difference(Exports))
```

```
caf |> autoplot(d.exports) + ggtitle("Differenced Exports")
```

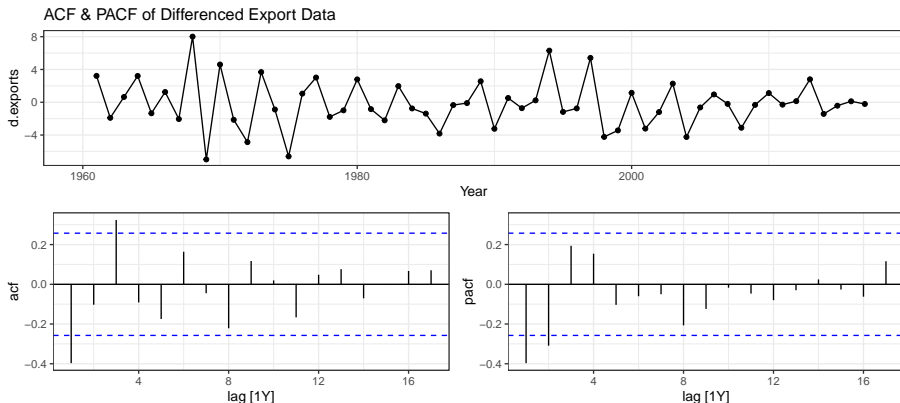


```
caf |> features(d.exports, unitroot_kpss)
```

```
## # A tibble: 1 x 3
```

```
##   Country                kpss_stat kpss_pvalue  
##   <fct>                  <dbl>         <dbl>  
## 1 Central African Republic 0.0922         0.1
```

```
caf |> gg_tsdisplay(d.exports, plot_type = "partial") +  
  labs(title = "ACF & PACF of Differenced Export Data")
```



Candidate Models:

- ARIMA(1,1,0)
- ARIMA(2,1,0)
- ARIMA(0,1,1)
- ARIMA(0,1,3)
- ARIMA(1,1,1)
- ARIMA(1,1,3)
- ARIMA(2,1,1)
- ARIMA(2,1,3)




```
caf.fits <- caf |> model(  
  arima110 = ARIMA(Exports ~ pdq(1,1,0)),  
  arima210 = ARIMA(Exports ~ pdq(2,1,0)),  
  arima011 = ARIMA(Exports ~ pdq(0,1,1)),  
  arima013 = ARIMA(Exports ~ pdq(0,1,3)),  
  arima111 = ARIMA(Exports ~ pdq(1,1,1)),  
  arima113 = ARIMA(Exports ~ pdq(1,1,3)),  
  arima211 = ARIMA(Exports ~ pdq(2,1,1)),  
  arima213 = ARIMA(Exports ~ pdq(2,1,3))  
)
```

Select “Best” Model

```
caf.fits |> glance() |> select(.model:BIC) |>  
  arrange(AICc)
```

```
## # A tibble: 7 x 6  
##   .model    sigma2 log_lik    AIC    AICc    BIC  
##   <chr>      <dbl>   <dbl> <dbl> <dbl> <dbl>  
## 1 arima210    6.71   -134.  275.  275.  281.  
## 2 arima013    6.54   -133.  274.  275.  282.  
## 3 arima011    7.01   -136.  276.  276.  280.  
## 4 arima211    6.69   -134.  275.  276.  284.  
## 5 arima113    6.60   -133.  276.  277.  286.  
## 6 arima110    7.22   -137.  278.  278.  282.  
## 7 arima111    7.07   -136.  277.  278.  284.
```

```
caf.fits |>  
  select(arima210) |>  
  gg_tsresiduals()
```




```
caf.fits |> augment() |>  
  filter(.model == "arima210") |>  
  features(.innov, ljung_box, lag = 10, df = 2)
```

```
## # A tibble: 1 x 4
```

##	Country	.model	lb_stat	lb_pvalue
##	<fct>	<chr>	<dbl>	<dbl>
## 1	Central African Republic	arima210	10.7	0.382

```
caf.fits |>  
  forecast(h = 10) |>  
  filter(.model == "arima210") |>  
  autoplot(caf)
```



Section 6

Seasonal ARIMA Models

Seasonal ARIMA Models

- Seasonal pattern in a year
 - Seasonal pattern year after year
 - Seasonal differences for nonstationary data
- The notation for a seasonal ARIMA model:

$$\text{ARIMA} \quad \underbrace{(p, d, q)}_{\text{Nonseasonal part}} \quad \underbrace{(P, D, Q)_m}_{\text{Seasonal part}}$$

where

- p is regular AR terms
- d is regular differences
- q is regular MA terms
- m is the seasonal period (e.g., number of observations per year)
- P is seasonal autoregressive terms at lag m
- D is seasonal difference at lag m
- Q is seasonal MA terms at lag m .

ACF/PACF Pattern for Seasonal lags

The seasonal part of an AR or MA model will be seen in the **seasonal lags** of the PACF and ACF.

For example, an $\text{ARIMA}(0, 0, 0)(0, 0, 1)_{12}$ model will show:

- a spike at lag 12 in the ACF but no other significant spikes;
- exponential decay in the seasonal lags of the PACF (i.e., at lags 12, 24, 36, ...).

ACF/PACF Pattern for Seasonal lags

Similarly, an $\text{ARIMA}(0, 0, 0)(1, 0, 0)_{12}$ model will show:

- exponential decay in the seasonal lags of the ACF;
- a single significant spike at lag 12 in the PACF.

When considering the appropriate seasonal orders for a seasonal ARIMA model, we restrict attention to the **seasonal lags**.


The modelling procedure is almost the same as for non-seasonal data, except that we need to select seasonal AR and MA terms as well as the non-seasonal components of the model.

Seasonal ARIMA model in Motion

```
leisure <- us_employment |>
  filter(Title == "Leisure and Hospitality", year(Month) > 2000) |>
  mutate(Employed = Employed/1000) |>
  select(Month, Employed)

autoplot(leisure, Employed) +
  labs(title = "US employment: leisure and hospitality",
       y="Number of people (millions)")
```


US employment: leisure and hospitality



The data are non-stationary with trend and seasonality.

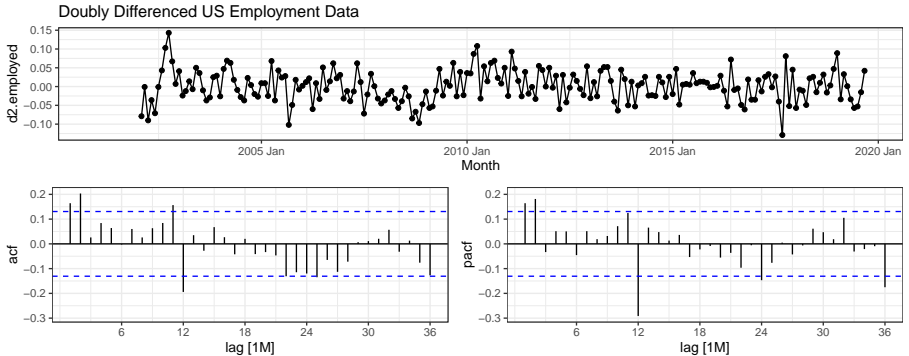
Let us first seasonally difference the data

```
leisure |> mutate(d.employed = difference(Employed, 12)) |>  
  gg_tsdisplay(d.employed, plot_type = "partial") +  
  labs(title = "Seasonally Differenced US Employment Data")
```



The series still appears to be nonstationary with a non-linear trend. Take further first difference

```
leisure |> mutate(d2.employed = difference(Employed, 12) |>  
  difference()) |>  
gg_tsdisplay(d2.employed, plot_type = "partial", lag_max = 36) +  
labs(title = "Doubly Differenced US Employment Data")
```



Candidate Models

- ➊ ARIMA(0, 1, 2)(1, 1, 0)₁₂
- ➋ ARIMA(0, 1, 2)(0, 1, 1)₁₂
- ➌ ARIMA(0, 1, 2)(1, 1, 1)₁₂
- ➍ ARIMA(2, 1, 0)(1, 1, 0)₁₂
- ➎ ARIMA(2, 1, 0)(0, 1, 1)₁₂
- ➏ ARIMA(2, 1, 0)(1, 1, 1)₁₂
- ➐ ARIMA(2, 1, 2)(1, 1, 0)₁₂
- ➑ ARIMA(2, 1, 2)(0, 1, 1)₁₂
- ➒ ARIMA(2, 1, 2)(1, 1, 1)₁₂

```
s.arima.fit <- leisure |> model(  
  arima012110 = ARIMA(Employed ~ pdq(0,1,2) + PDQ(1,1,0)),  
  arima012011 = ARIMA(Employed ~ pdq(0,1,2) + PDQ(0,1,1)),  
  arima012111 = ARIMA(Employed ~ pdq(0,1,2) + PDQ(1,1,1)),  
  arima210110 = ARIMA(Employed ~ pdq(2,1,0) + PDQ(1,1,0)),  
  arima210011 = ARIMA(Employed ~ pdq(2,1,0) + PDQ(0,1,1)),  
  arima210111 = ARIMA(Employed ~ pdq(2,1,0) + PDQ(1,1,1)),  
  arima212110 = ARIMA(Employed ~ pdq(2,1,2) + PDQ(1,1,0)),  
  arima212011 = ARIMA(Employed ~ pdq(2,1,2) + PDQ(0,1,1)),  
  arima212111 = ARIMA(Employed ~ pdq(2,1,2) + PDQ(1,1,1))  
)
```

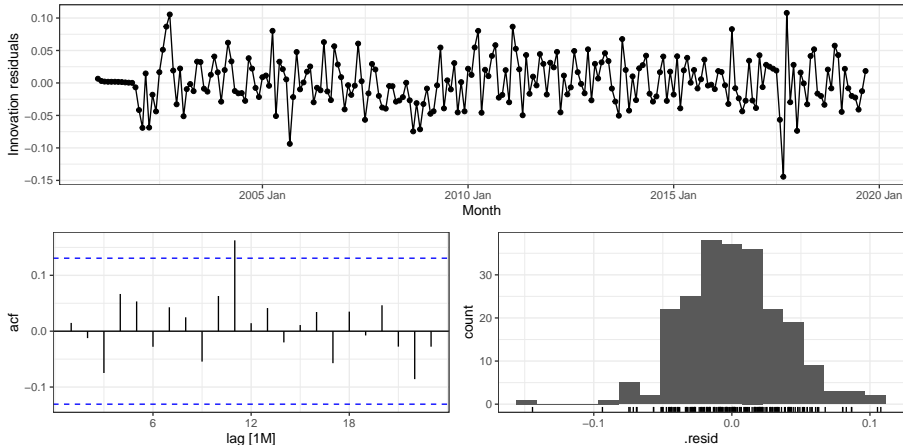
```
s.arima.fit |> glance() |>
  arrange(AICc) |> select(.model:BIC)
```

```
## # A tibble: 9 x 6
```

##	.model	sigma2	log_lik	AIC	AICc	BIC
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	arima210111	0.00142	395.	-780.	-780.	-763.
## 2	arima012111	0.00142	395.	-779.	-779.	-763.
## 3	arima210011	0.00145	392.	-776.	-776.	-763.
## 4	arima212111	0.00143	395.	-776.	-775.	-752.
## 5	arima012011	0.00146	391.	-775.	-775.	-761.
## 6	arima212011	0.00146	393.	-773.	-773.	-753.
## 7	arima210110	0.00154	387.	-766.	-765.	-752.
## 8	arima012110	0.00156	386.	-764.	-764.	-751.
## 9	arima212110	0.00156	387.	-762.	-762.	-742.

Check Model Residuals

```
s.arima.fit |> select(arima210111) |> gg_tsresiduals()
```



Lag 11 might be concerning. Conduct a Ljung-Box test.

Check Model Residuals

```
s.arima.fit |> select(arima210111) |>  
  augment() |> features(.innov, lbjung_box, lag = 24,  
                        dof = 4)
```

```
## # A tibble: 1 x 3  
##   .model      lb_stat lb_pvalue  
##   <chr>      <dbl>    <dbl>  
## 1 arima210111    16.6      0.680
```

False Alarm: We actually fail to reject null of no autocorrelation. Our residuals are indeed white noise.

Forecasting with Preferred Model

```
s.arima.fit |> forecast(h = "3years") |>  
  filter(.model == "arima210111") |>  
  autoplot(leisure) +  
  labs(title = "Forecasted US Employment: Leisure & Hospitality",  
        subtitle = bquote("ARIMA(2,1,0)(1,1,1)"[12]))
```

Forecasted US Employment: Leisure & Hospitality

ARIMA(2,1,0)(1,1,1)₁₂

